




A · Stacking hard questions SATURATES**B · Modality diversity is ~inert****C · Discretize → cleaner estimate**

Mini V-L-A, RT-1/fractal, SUBN=4000, 3 seeds. (A) benefit of a low-recoverability obs-mix rises 1→2 then plateaus — it is the PRESENCE of a hard member, not the count, that sets the mix (composition law). (B) swapping the second member across modalities (obs/action/pose) barely moves the benefit → the lever is recoverability, not modality diversity. (C) the next action is a genuinely low-recoverability target; a discrete K=16 cross-entropy estimate is sign-stable and 3.4× lower-variance than continuous R^2 — why FAST/ $\pi 0.5$'s discrete action tokens are a cleaner low-recoverability signal.